

USC CSCI 567 HOMEWORK 3 SOLUTIONS

THAMMEGOWDA NARAYANASWAMY

tnarayan@usc.edu

ID : 2074-6694-39

Department of Computer Science

Viterbi School of Engineering

University of Southern California

Los Angeles, CA

October 17, 2016

1. BIAS VARIANCE TRADE-OFF

a. Closed form Solution for Ridge Regression

Given :

$$\hat{\beta}_\lambda = \operatorname{argmin}_\beta \left\{ \frac{1}{n} \sum_{i=1}^n n(y_i - x_i^T \beta)^2 + \lambda \| \beta \|_2^2 \right\}$$

In matrix form:

$$\hat{\beta}_\lambda = \operatorname{argmin}_\beta \left\{ \frac{1}{n} (Y - X\beta)^T (Y - X\beta) + \lambda \beta^T \beta \right\}$$

where $X \in R^{n \times p}$, $Y \in R^{n \times 1}$ and $\beta \in R^{p \times 1}$. We know that the above expression (assuming it is convex) is minimum when its derivative is zero.

$$\begin{aligned} \frac{\partial}{\partial \beta} \left[\frac{1}{n} (Y - X\beta)^T (Y - X\beta) + \lambda \beta^T \beta \right] &= 0 \\ \frac{1}{n} \frac{\partial}{\partial \beta} \left[Y^T Y - Y^T X\beta - (X\beta)^T Y + (X\beta)^T (X\beta) + \lambda \beta^T \beta \right] &= 0 \\ \frac{\partial}{\partial \beta} \left[-Y^T X\beta - \beta^T X^T Y + \beta^T X^T X\beta + \lambda \beta^T \beta \right] &= 0 \\ -Y^T X - X^T Y + 2\beta^T X^T X + 2\lambda \beta &= 0 \\ -2X^T Y + 2X^T X\beta + 2\lambda \beta &= 0 \\ [X^T X + \lambda I]\beta &= X^T Y \end{aligned}$$

When $X^T X + \lambda I$ is invertible,

$$\begin{aligned} \beta &= [X^T X + \lambda I]^{-1} X^T Y \\ \Rightarrow \hat{\beta}_\lambda &= [X^T X + \lambda I]^{-1} X^T Y \end{aligned}$$

We are told that the target label y is from linear model plus an additional Gaussian noise ϵ .

\Rightarrow Distribution of ϵ is $N(0, \sigma^2)$

\Rightarrow By applying affine transformations, the distribution of $Y \approx N(X\beta^*, \sigma^2)$

\Rightarrow Since $\hat{\beta}_\lambda = H_\lambda Y$, where $H_\lambda = [X^T X + \lambda I]^{-1} X^T$

The distribution of $\hat{\beta}_\lambda \approx N(H_\lambda X\beta^*, H_\lambda^T \sigma^2 I H_\lambda)$

b. Bias Term

$$\begin{aligned} \text{bias} &= \mathbb{E}[x^T \hat{\beta}_\lambda] - x^T \beta^* \\ &= x^T [\mathbb{E}[\hat{\beta}_\lambda] - \beta^*] \\ &= x^T [H_\lambda X\beta^* - \beta^*] \quad \because \text{from part a, the mean of Gaussian} \\ &\quad \text{where } H_\lambda = [X^T X + \lambda I]^{-1} X^T \\ \text{bias} &= x^T (H_\lambda X - I) \beta^* \end{aligned}$$

c. Variance Term

$$\begin{aligned}
 \text{variance} &= \mathbb{E} \left[(x^T \hat{\beta}_\lambda - \mathbb{E}[x^T \hat{\beta}_\lambda])^2 \right] \\
 &= \mathbb{E} \left[(x^T (\hat{\beta}_\lambda - \mathbb{E}[\hat{\beta}_\lambda]))^2 \right] \\
 &= \left(x^T (\mathbb{E}[\hat{\beta}_\lambda - \mathbb{E}[\hat{\beta}_\lambda]]) \right)^2 \\
 &= x^T (\mathbb{E}[\hat{\beta}_\lambda - \mathbb{E}[\hat{\beta}_\lambda]])^2 x \\
 &= x^T \text{Var}(\hat{\beta}_\lambda) x \\
 \text{Variance} &= x^T H_\lambda^T \sigma^2 I H_\lambda x & \because \text{From Part - a} \\
 \text{Where, } H_\lambda &= [X^T X + \lambda I]^{-1} X^T
 \end{aligned}$$

d. Impact of λ on Squared Error

From the bias-variance theorem we have, squared error is proportional to sum of *variance* and *bias*²

The impact of λ on squared error is as follows:

- When λ is smaller, the bias is small but variance is high. Thus the model tight fits to seen samples with less generalization.
- When λ is large, the bias is larger but variance is smaller. Thus the model under fits the samples with poor prediction.
- Thus we cannot have least bias and least variance at the same time, we need to settle for a trade off between these.

2. KERNEL CONSTRUCTION

a. Positive linear combination of valid kernels is a valid kernel

To prove: $k_3(x, x') = a_1 k_1(x, x') + a_2 k_2(x, x')$ is a valid kernel

Given: k_1 and k_2 are positive definite kernel functions. $a_1, a_2 \geq 0$

Proof: Since k_1 and k_2 are positive definite kernel functions, we have from Mercer theorem:

$$\int_x \int_{x'} k_1(x, x') f(x) f(x') dx dx' \geq 0 \quad (1)$$

And

$$\int_x \int_{x'} k_2(x, x') f(x) f(x') dx dx' \geq 0 \quad (2)$$

$$\implies k_3(x, x') = a_1 \int_x \int_{x'} k_1(x, x') f(x) f(x') dx dx' + a_2 \int_x \int_{x'} k_2(x, x') f(x) f(x') dx dx'$$

From equations 1, 2, and $a_1 \geq 0, a_2 \geq 0$

\implies R.H.S is ≥ 0

$\implies k_3$ is a positive definite kernel function

b.

To prove: $k_4(x, x') = f(x)f(x')$ is a valid kernel

Given: $f(\cdot)$ is an arbitrary real valued function

Proof:

Let x and $x' \in \mathbb{R}^n$

We can represent $k_4(x, x')$ as matrix K_4 , given by

$$K_4 = \begin{bmatrix} f(x_1)f(x'_1) & f(x_1)f(x'_2) & \dots & f(x_1)f(x'_n) \\ \dots & \dots & \dots & \dots \\ f(x_n)f(x'_1) & f(x_n)f(x'_2) & \dots & f(x_n)f(x'_n) \end{bmatrix}$$

The matrix can be decomposed into product of two vectors F^T and F , $F \in \mathbb{R}^n$, such that $K_4 = F^T \cdot F$ where

$$F = \begin{bmatrix} f(x_1) \\ f(x_2) \\ \dots \\ f(x_n) \end{bmatrix}$$

To prove that k_4 is a valid kernel, we have to prove that the matrix K_4 is positive definite.

Let us consider, $y \in \mathbb{R}^n$

$$\begin{aligned} y^T K_4 y &= y^T F F^T y \\ &= (F^T y)^T F^T y \\ &= \| F^T y \|_2^2 \end{aligned}$$

Since L_2 norm is always positive, we can infer that for any $y \in \mathbb{R}^n$, $y^T K_4 y \geq 0$
 $\implies k_4$ is a valid kernel by definition

c. The product of valid Kernel is a valid kernel

To prove: $k_5(x, x') = k_1(x, x')k_2(x, x')$ is a valid kernel

Given: k_1 and k_2 are positive definite kernel functions.

Proof:

Let us consider Kernel matrices K_1 , K_2 , and K_5 corresponding to k_1 , k_2 , and k_5 . By definition of kernel functions we can say that K_1 and K_2 are:

- Symmetric
- Positive semi definite

We know that, product of two symmetric positive semi definite matrices is also symmetric positive semidefinite matrix. (necessary and sufficient for valid kernel)

\implies Thus k_5 is a valid kernel.

3. KERNEL REGRESSION

a. Optimal value of w

Given: the cost function

$$J(w) = \min_w \sum_n (y_i - w^T x_i)^2 + \lambda \|w\|_2^2$$

Where $\lambda \geq 0$ is regularization coefficient.

We know that, J is minimal when $\frac{\partial J}{\partial w}$

$$\begin{aligned} \frac{\partial}{\partial w} \left[\frac{1}{n} (y - Xw)^T (y - Xw) + \lambda w^T w \right] &= 0 \\ \frac{\partial}{\partial w} \left[y^T y - y^T Xw - (Xw)^T y + (Xw)^T (Xw) + \lambda w^T w \right] &= 0 \\ \frac{\partial}{\partial w} \left[-y^T Xw - w^T X^T y + w^T X^T Xw + \lambda w^T w \right] &= 0 \\ -y^T X - X^T y + 2w^T X^T X + 2\lambda w &= 0 \\ -2X^T y + 2X^T Xw + 2\lambda w &= 0 \\ [X^T X + \lambda I]w &= X^T y \end{aligned}$$

When $X^T X + \lambda I_D$ is invertible,

$$\implies w^* = [X^T X + \lambda I_D]^{-1} X^T y$$

b.

Given that $x_i \in \mathbb{R}^D$ has been transformed to $\phi_i(x) \in \mathbb{R}^T$

Thus the input matrix is $\Phi \in \mathbb{R}^{N \times T}$. Substituting the input matrix in the equation from part a, we get

$$w^* = [\Phi^T \Phi + \lambda I_T]^{-1} \Phi^T y \quad (3)$$

Let us consider the expression,

$$\begin{aligned} [\Phi^T \Phi + \lambda I_D] \Phi^T &= [\Phi^T \Phi \Phi^T + \lambda \Phi^T] && \because \text{Distributing } \Phi^T \\ [\Phi^T \Phi + \lambda I_D] \Phi^T &= \Phi^T [\Phi \Phi^T + \lambda I_N] && \because \text{Moving to the left side } \Phi^T \\ \Phi^T &= [\Phi^T \Phi + \lambda I_D]^{-1} \Phi^T [\Phi \Phi^T + \lambda I_N] && \because \text{Inverse Multiplied on left} \\ \Phi^T [\Phi \Phi^T + \lambda I_N]^{-1} &= [\Phi^T \Phi + \lambda I_D]^{-1} \Phi^T && \because \text{Inverse Multiplied on Right} \end{aligned}$$

Substituting the result in 3, we get

$$w^* = \Phi^T [\Phi \Phi^T + \lambda I_N]^{-1} y \quad (4)$$

c.

Given a testing sample $\phi(x)$, the prediction is given by,

$$\begin{aligned}\hat{y} &= w^{*T} \phi(x) \\ &= (\Phi^T [\Phi \Phi^T + \lambda I_N]^{-1} y)^T \phi(x) \\ &= (y^T [(\Phi \Phi^T + \lambda I_N)^T]^{-1} \Phi \phi(x)) \\ &= (y^T [(\Phi \Phi^T)^T + \lambda I_N]^{-1} \Phi \phi(x))\end{aligned}$$

Lets define a matrix $K = \Phi \Phi^T \in \mathbb{R}^{N \times N}$,

and $k(x) = \Phi \phi(x) \in \mathbb{R}^T$

Since K is symmetric, $K^T = K$

$$\therefore \hat{y} = (y^T [K + \lambda I_N]^{-1} k(x))$$

d.

When the dimension of feature transformation $\phi(x) \in \mathbb{R}^T, T$, is large, kernel regressions are computationally efficient than linear ridge regression. \therefore kernel regression computes the kernel matrix K which is $\mathbb{R}^{N \times N}$ without actually computing the transformations $\phi(x)$. Without kernel trick, we have to compute $T \times T$ matrix which is an expensive operation when $T \gg N$.

- **Learning:**

Linear Ridge Regression: $O(T^3)$ time and $O(T^2)$ memory; T is the dimension of new feature space.

Kernel Regression: $O(N^3)$ time and $O(N^2)$ memory; N is the number of examples in training set.

- **Prediction:**

Linear Ridge Regression: $O(T)$ time, T is the dimension of new feature space.

Kernel Regression: $O(D)$ time, D is the dimensions of original feature space

4. SUPPORT VECTOR MACHINES

1. Linearly separable?


No. They are not linearly separable (in 2d euclidean space).

2. Maximum margin decision surface

$$w = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$

3. Make it non separable

The data can be easily made non separable by adding a negative label at the same as a positive example (or vice versa). In practice, this issue is called training data contamination.



Example: Adding positive example at (-1, 1)


4. Kernel name

$$\begin{aligned}
 K(x, x) &= \langle \phi(x), \phi(x) \rangle = \phi(x)^T \phi(x) \\
 &= [1 \ x_1 \ x_2 \ x_1 x_2]^T [1 \ x_1 \ x_2 \ x_1 x_2] \\
 &= 1 + x_1 + x_2 + x_1 x_2 + x_1 + x_1^2 + x_1 x_2 + x_1^2 x_2 + x_2 + x_1 x_2 + x_2^2 + x_1 x_2^2 \\
 &\quad + x_1 x_2 + x_1^2 x_2 + x_1 x_2^2 + x_1^2 x_2^2
 \end{aligned}$$

This is a polynomial kernel.


5. SVMs AND THE SLACK PENALTY C

1.



When $C \rightarrow \inf$, SVM tends to minimize misclassification errors, even if the margin between the decision boundary plane and the training examples is small.

2.

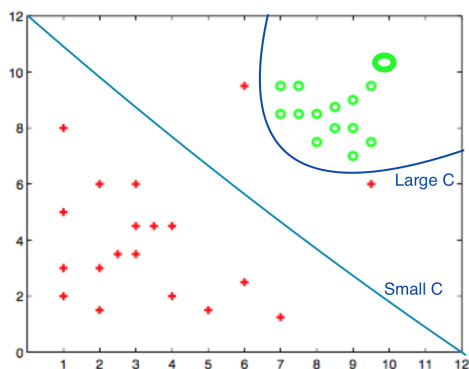


When $C \approx 0$, SVM tends to maximize the margin between the decision boundary and the training examples. This might even result in few misclassification errors.

3.

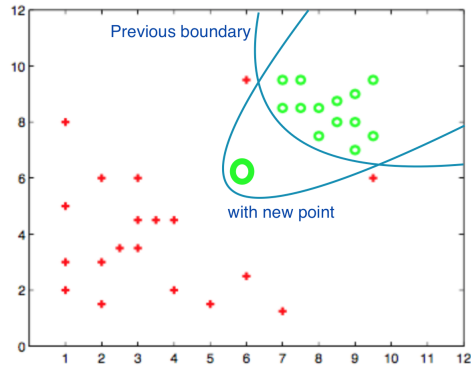
The precision of Red Plus points and the Recall of the Green Circle points will be high. Depending on the value of slack penalty (as discussed in the previous two sections), the recall of the Red Plus points and the precision of the Green circle points will vary and are most likely to perform poorly. This is because there are few red plus points closer to green circle points which affect the position of the decision boundary.

4.



A point is added at the top right region of the plane. Since it is far from being a 'support vector', it doesn't change the decision boundary.

5.



A point is added in between the decision boundary region. Since it becomes 'support vector', it participates in forming the decision boundary.

PROGRAMMING

Bias Variance

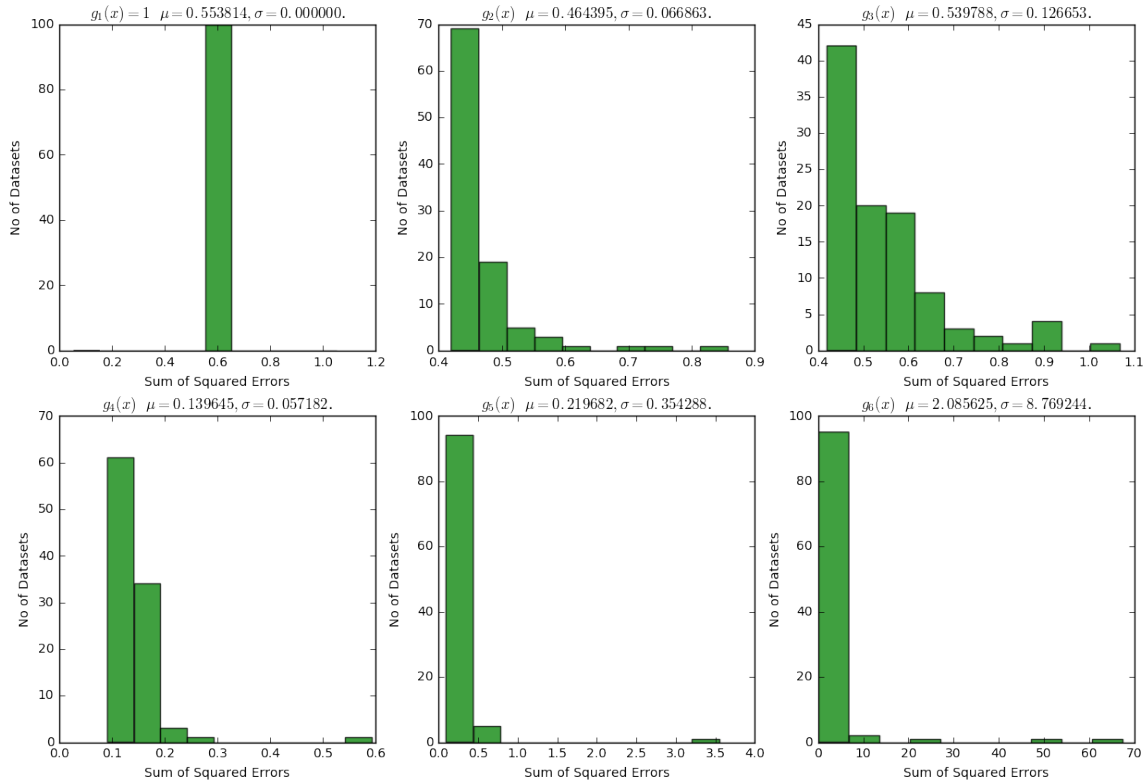
a. 10 samples per dataset

Note: The evaluation is performed as follows based on the method suggested by TA's in one of the discussion sessions:

- Generated 100 datasets with 10 points each, $D_1, D_2, D_3 \dots D_{100}$
- Built a model $h_{D_i}(\cdot)$ for each 100 datasets.
- A news dataset D_t is generated for testing.
- The bias and variance are measured based on the performance of 100 models on the D_t dataset.

Histograms

The histograms are as follows

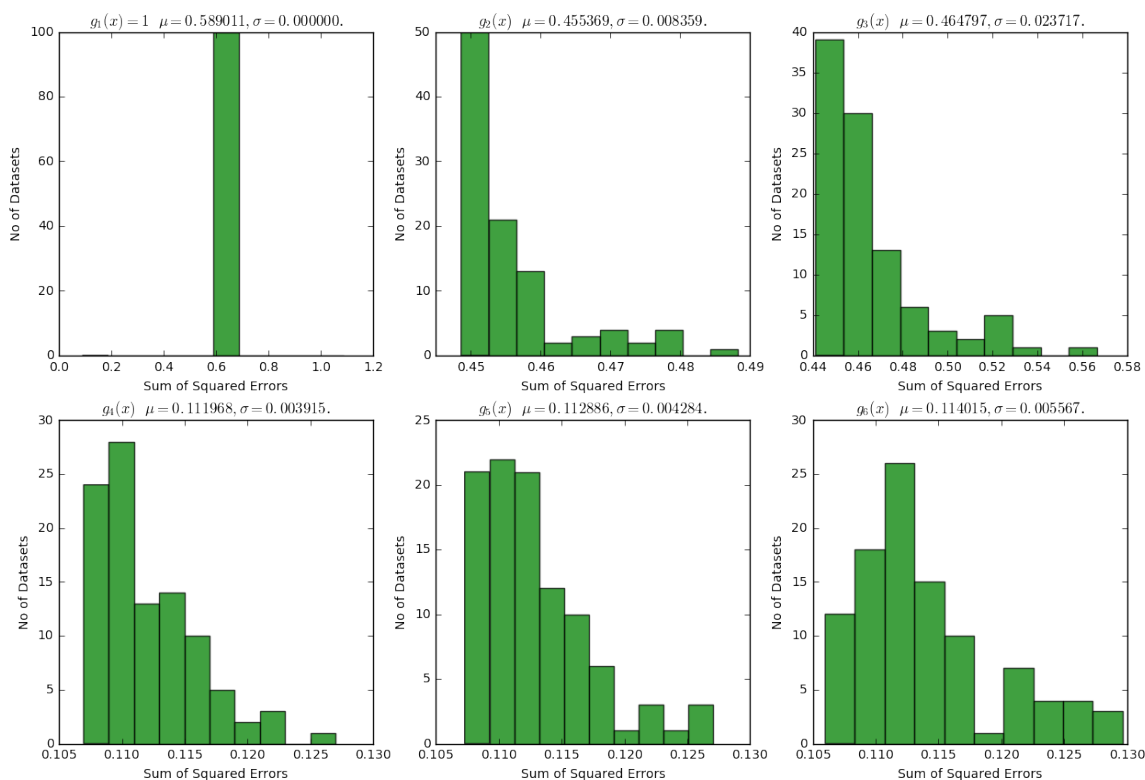


The recorded $bias^2$ and variance are as follows:

Function	Variance	$Bias^2$
g1	0.00000000	0.55381366
g2	0.04447961	0.41991505
g3	0.11798539	0.42180300
g4	0.04600876	0.09363664
g5	0.12355812	0.09612371
g6	1.98816964	0.09745515

b. 100 Samples per dataset

The histograms are as follows



The recorded $bias^2$ and variance are as follows:

Function	Variance	$Bias^2$
g1	0.00000000	0.58901086
g2	0.00464569	0.45072357
g3	0.01270188	0.45209526
g4	0.00269852	0.10926987
g5	0.00366258	0.10922298
g6	0.00511599	0.10889857

c.

Key observations from this exercise:

- As the model complexity increased, the bias decreased and the variance increased..
- When we added more number of samples to the dataset, we saw that the bias and variance continue to vary (bias decreased, variance increased) with increase in model complexity, however the magnitude of change is less.

d. Ridge Regression

Lambda	Variance	$Bias^2$
0.001	0.00291082	0.110499
0.003	0.00290983	0.110496
0.01	0.00290636	0.110483
0.03	0.00289664	0.110451
0.1	0.00286453	0.110368
0.3	0.00278789	0.110376
1	0.00265098	0.11267

Analysis: As λ increased, the bias increased and the variance decreased. Thus λ parameter can be used to control tight fitting by introducing the bias.

Support Vector Machines**libSVM setup**

Using anaconda distribution:

```
conda install -c conda-forge libsvm=3.21
```

This installs two command line tools called `svm_train` `svm_predict`

Using the Linear Kernel

The following results were recorded by using linear kernel type. The unit of average time is seconds for all the following results.

```
# LINEAR KERNEL
#           C AvgTime  Accuracy%
1  0.000244141 0.213693 55.75
2  0.000976562 0.201526 55.75
3  0.00390625 0.197195 55.75
4  0.015625 0.190556 55.75
5  0.0625 0.196112 74.90
6  0.25 0.147726 92.05
7  1 0.100852 92.80
8  4 0.065552 95.10
9  16 0.053958 96.40
```

Observation: The accuracy increased as we increased the slack penalty C , while the average training time decreased.

Using the polynomial and RBF kernels

The following results were recorded:

```
# POLYNOMIAL KERNEL
#           C Degree AvgTime  Accuracy%
```

1	0.015625	1	0.193540	55.75	
2	0.015625	2	0.201199	55.75	
3	0.015625	3	0.200625	55.75	
4	0.0625	1	0.174386	90.35	
5	0.0625	2	0.198155	89.00	
6	0.0625	3	0.198131	74.90	
7	0.25	1	0.103184	91.15	
8	0.25	2	0.116746	92.15	
9	0.25	3	0.147319	92.05	
10		1	1	0.073813	92.85
11		1	2	0.079659	93.30
12		1	3	0.100888	92.80
13		4	1	0.056519	94.35
14		4	2	0.063589	94.90
15		4	3	0.067277	95.10
16		16	1	0.054508	94.30
17		16	2	0.057410	95.95
18		16	3	0.055024	96.40
19		64	1	0.056362	94.15
20		64	2	0.050879	96.65
21		64	3	0.051658	97.00
22		256	1	0.068880	94.50
23		256	2	0.053696	96.80
24		256	3	0.049736	96.90
25		1024	1	0.154141	94.50
26		1024	2	0.056765	96.50
27		1024	3	0.053674	96.55
28		4096	1	0.381076	94.50
29		4096	2	0.059383	96.75
30		4096	3	0.056649	96.60
31		16384	1	1.563308	94.70
32		16384	2	0.065256	96.45
33		16384	3	0.055341	96.60

RBF KERNEL

#	C	Gamma	AvgTime	Accuracy%
1	0.015625	0.000061	0.196857	55.75
2	0.015625	0.000244	0.199696	55.75
3	0.015625	0.000977	0.193203	55.75
4	0.015625	0.003906	0.201737	55.75
5	0.015625	0.015625	0.200773	56.00
6	0.015625	0.062500	0.198204	87.30
7	0.015625	0.250000	0.200537	60.30
8	0.0625	0.000061	0.199040	55.75
9	0.0625	0.000244	0.200937	55.75

10	0.0625	0.000977	0.196819	55.75
11	0.0625	0.003906	0.203154	64.40
12	0.0625	0.015625	0.165145	90.55
13	0.0625	0.062500	0.128067	92.20
14	0.0625	0.250000	0.172642	92.05
15	0.25	0.000061	0.200721	55.75
16	0.25	0.000244	0.199568	55.75
17	0.25	0.000977	0.199016	66.60
18	0.25	0.003906	0.154790	90.85
19	0.25	0.015625	0.103935	91.40
20	0.25	0.062500	0.082446	93.10
21	0.25	0.250000	0.115713	95.55
22	1	0.000061	0.196611	55.75
23	1	0.000244	0.198686	67.85
24	1	0.000977	0.152969	91.05
25	1	0.003906	0.099703	91.40
26	1	0.015625	0.074173	93.50
27	1	0.062500	0.063382	95.90
28	1	0.250000	0.085199	97.15
29	4	0.000061	0.197674	67.95
30	4	0.000244	0.151894	90.85
31	4	0.000977	0.102306	91.45
32	4	0.003906	0.072589	93.35
33	4	0.015625	0.056027	94.90
34	4	0.062500	0.054397	96.80
35	4	0.250000	0.074961	96.95
36	16	0.000061	0.154246	90.85
37	16	0.000244	0.097384	91.50
38	16	0.000977	0.074993	93.45
39	16	0.003906	0.057111	94.55
40	16	0.015625	0.053216	95.85
41	16	0.062500	0.052186	97.10
42	16	0.250000	0.075502	96.70
43	64	0.000061	0.097574	91.50
44	64	0.000244	0.069576	93.45
45	64	0.000977	0.058329	94.45
46	64	0.003906	0.056718	94.50
47	64	0.015625	0.054045	96.75
48	64	0.062500	0.049619	97.00
49	64	0.250000	0.074851	96.70
50	256	0.000061	0.071727	93.45
51	256	0.000244	0.058747	94.40
52	256	0.000977	0.056307	94.30
53	256	0.003906	0.055413	95.90
54	256	0.015625	0.054942	97.20

55	256	0.062500	0.047539	96.85
56	256	0.250000	0.074552	96.70
57	1024	0.000061	0.057464	94.40
58	1024	0.000244	0.057587	94.40
59	1024	0.000977	0.056671	94.60
60	1024	0.003906	0.066289	96.70
61	1024	0.015625	0.057694	96.90
62	1024	0.062500	0.049985	96.85
63	1024	0.250000	0.076784	96.70
64	4096	0.000061	0.055695	94.40
65	4096	0.000244	0.058861	94.35
66	4096	0.000977	0.073948	96.20
67	4096	0.003906	0.086435	97.10
68	4096	0.015625	0.065810	96.95
69	4096	0.062500	0.050349	96.85
70	4096	0.250000	0.074639	96.70
71	16384	0.000061	0.060005	94.25
72	16384	0.000244	0.080497	94.70
73	16384	0.000977	0.114950	96.55
74	16384	0.003906	0.110726	96.75
75	16384	0.015625	0.065220	96.95
76	16384	0.062500	0.048876	96.85
77	16384	0.250000	0.073846	96.70

Observation : The best cross validation accuracy is 97.20% and the corresponding parameters are $Kernel = RBF$; $C = 256$; $gamma = 0.015625$

The model built from these parameters recorded 98.5% accuracy on the test dataset.